

Keane Chan Jun Yu
A0205678W

GitHub Repo: <https://github.com/keanecjy/CS3219-TaskA3>

CS3219-TaskA3 - Ingress Controller

Deployment

Verify that NGINX Ingress controller is up

```
kubectl get pods -n ingress-nginx
```

```
$ kubectl get pods -n ingress-nginx
```

NAME	READY	STATUS	RESTARTS	AGE
ingress-nginx-admission-create--1-9k6gb	0/1	Completed	0	43h
ingress-nginx-admission-patch--1-wkqcd	0/1	Completed	1	43h
ingress-nginx-controller-69bdbc4d57-c8wn6	1/1	Running	1 (82m ago)	43h

Deploy and create service

```
kubectl apply -f deployment.yml
```

```
kubectl get deployments
```

```
$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
ingress-demo	2/2	2	2	112m

```
kubectl get service
```

```
$ kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
ingress-demo	NodePort	10.98.77.8	<none>	80:32334/TCP	83m

Add Ingress controller. It can take awhile for the ip address to be set.

```
kubectl apply -f ingress.yml  
kubectl get ingress
```

```
$ kubectl get ingress
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
ingress-demo	nginx	ingress.demo.io	192.168.49.2	80	18m

Using this ip address, go to /etc/hosts file and add `<IP_ADDRESS> ingress.demo.io` at the bottom of the file.

Now we can access the endpoint at <http://ingress.demo.io>

```
curl -i http://ingress.demo.io
```

```
$ curl -i http://ingress.demo.io  
HTTP/1.1 200 OK  
Date: Sat, 09 Oct 2021 14:14:57 GMT  
Content-Type: text/plain; charset=utf-8  
Content-Length: 69  
Connection: keep-alive  
Set-Cookie: INGRESSCOOKIE=1633788898.124.173.161235; Expires=Mon, 11-Oct-21 14:14:57 GMT; Max-Age=172800; Path=/; HttpOnly  
  
Hello, world!  
Version: 1.0.0  
Hostname: ingress-demo-55974b588f-9nkw4
```

As seen here, we see that there is a cookie set in the field. Given the same cookie, the routing will always direct back to the same webpage.

CS3219-TaskA3 - Horizontal pod auto-scaler

Create deployment, expose it and attach service

```
kubectl apply -f hpa.yml
```

```
$ kubectl apply -f hpa.yml  
deployment.apps/hpa-demo created  
service/hpa-demo created
```

Check if deployment exist

```
kubectl get deployments
```

```
$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hpa-demo	1/1	1	1	23s

Create hpa

```
kubectl autoscale deployment hpa-demo --cpu-percent=50 --min=1 --max=10
```

```
$ kubectl autoscale deployment hpa-demo --cpu-percent=50 --min=1 --max=10
horizontalpodautoscaler.autoscaling/hpa-demo autoscaled
```

Check status of autoscaler (might take some time to show %). We should expect 0% since we are not sending any requests to the server.

```
kubectl get hpa
```

```
$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo	Deployment/hpa-demo	1%/50%	1	10	1	50s

Increase load on the service

```
kubectl run -i --tty load-generator --rm --image=busybox --restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://hpa-demo; done"
```

```
$ kubectl run -i --tty load-generator --rm --image=busybox --restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://hpa-demo; done"
If you don't see a command prompt, try pressing enter.
OK!OK!OK!OK!OK!OK!
```

Open up a new terminal. After 1 minute, we should see higher cpu load:

```
kubectl get hpa
```

```
$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo	Deployment/hpa-demo	52%/50%	1	10	5	2m23s

We should also see more replicas created for the deployment

```
kubectl get deployment hpa-demo
```

```
$ kubectl get deployment hpa-demo
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
hpa-demo      5/5      5             5            3m47s
```

Now we stop the container with the busybox image using Ctrl-C (SIGINT). After 1 minute or so, cpu utilization will drop back to 0%

```
kubectl get hpa
```

```
$ kubectl get hpa
NAME          REFERENCE          TARGETS    MINPODS    MAXPODS    REPLICAS    AGE
hpa-demo      Deployment/hpa-demo  1%/50%    1          10         1           14m
```

After several minutes, the HPA will autoscale and decrease the no. of replicas down to 1.

```
kubectl get deployment hpa-demo
```

```
$ kubectl get deployment hpa-demo
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
hpa-demo      1/1      1             1            14m
```

Resources

1. <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/>